

Problem 2

(a) Check CIFAR10

```
In [1]: import torch
import torchvision
import torchvision.transforms as transforms
```

```
In [2]: transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

batch_size = 4

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified
Files already downloaded and verified

```
In [3]: trainset
```

```
Out[3]: Dataset CIFAR10
        Number of datapoints: 50000
        Root location: ./data
        Split: Train
        StandardTransform
        Transform: Compose(
            ToTensor()
            Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))
        )
```

```
In [4]: testset
```

```
Out[4]: Dataset CIFAR10
        Number of datapoints: 10000
        Root location: ./data
        Split: Test
        StandardTransform
        Transform: Compose(
            ToTensor()
            Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))
        )
```


```
In [5]: import matplotlib.pyplot as plt
import numpy as np

# functions to show an image

def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(trainloader)
images, labels = next(dataiter)


# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join(f'{classes[labels[j]]:5s}' for j in range(batch_size)))
```



bird deer plane frog

```
In [6]: # get some random training images
dataiter = iter(testloader)
images, labels = next(dataiter)

# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join(f'{classes[labels[j]]:5s}' for j in range(batch_size)))
```



cat ship ship plane

```
In [7]: images, labels = next(dataiter)
        images.shape # 4 is batch size
```

```
Out[7]: torch.Size([4, 3, 32, 32])
```

(b) Train 1 hidden layer ReLU

```
In [8]: import torch.nn as nn
        import torch.nn.functional as F

        class Net(nn.Module):
            def __init__(self, k=128):
                super(Net, self).__init__()
                self.k = k
                self.fc1 = nn.Linear(3 * 32 * 32, self.k, bias=False)
                self.fc2 = nn.Linear(self.k, 10, bias=False)

            def forward(self, x):
                x = x.view(-1, 3 * 32 * 32) # [batchsize, 3072]
                x = F.relu(self.fc1(x))
                x = self.fc2(x)
                return x

        model = Net(k=128)
```

```
In [9]: criterion = nn.MSELoss()
        optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
```

```
In [11]: # train the model
        num_epochs = 20
        train_loss_history = []
        train_acc_history = []
        test_loss_history = []
        test_acc_history = []

        # Loop through the number of epochs
        for epoch in range(num_epochs):
```

```
train_loss = 0.0
train_acc = 0.0
test_loss = 0.0
test_acc = 0.0

# set model to train mode
model.train()
# iterate over the training data
for inputs, labels in trainloader:
    optimizer.zero_grad()
    outputs = model(inputs)

    #compute the loss
    one_hot_labels = torch.nn.functional.one_hot(labels, num_classes=10)
    loss = criterion(outputs, one_hot_labels.float())
    loss.backward()
    optimizer.step()
    # increment the running loss and accuracy
    train_loss += loss.item()
    train_acc += (outputs.argmax(1) == labels).sum().item()

# calculate the average training loss and accuracy
train_loss /= len(trainloader)
train_loss_history.append(train_loss)
train_acc /= len(trainloader.dataset)
train_acc_history.append(train_acc)

# set the model to evaluation mode
model.eval()
with torch.no_grad():
    for inputs, labels in testloader:
        outputs = model(inputs)

        #compute the loss
        one_hot_labels = torch.nn.functional.one_hot(labels, num_classes=10)
        loss = criterion(outputs, one_hot_labels.float())

        test_loss += loss.item()
        test_acc += (outputs.argmax(1) == labels).sum().item()

# calculate the average validation loss and accuracy
test_loss /= len(testloader)
test_loss_history.append(test_loss)
test_acc /= len(testloader.dataset)
test_acc_history.append(test_acc)

print(f'Epoch {epoch+1}/{num_epochs}, train loss: {train_loss:.4f}, train acc: {train_acc:.4f}, test loss: {test_loss:.4f}, test acc: {test_acc:.4f}')
```

Epoch 1/20, train loss: 0.0944, train acc: 0.3199, val loss: 0.0879, val acc : 0.3390
Epoch 2/20, train loss: 0.0952, train acc: 0.3307, val loss: 0.1536, val acc : 0.2732
Epoch 3/20, train loss: 0.0942, train acc: 0.3431, val loss: 0.0938, val acc : 0.3313
Epoch 4/20, train loss: 0.0941, train acc: 0.3450, val loss: 0.0846, val acc : 0.3601
Epoch 5/20, train loss: 0.0935, train acc: 0.3508, val loss: 0.0916, val acc : 0.3296
Epoch 6/20, train loss: 0.0951, train acc: 0.3520, val loss: 0.0854, val acc : 0.3571
Epoch 7/20, train loss: 0.0928, train acc: 0.3552, val loss: 0.0904, val acc : 0.3580
Epoch 8/20, train loss: 0.0945, train acc: 0.3527, val loss: 0.0873, val acc : 0.3521
Epoch 9/20, train loss: 0.0949, train acc: 0.3622, val loss: 0.0998, val acc : 0.3059
Epoch 10/20, train loss: 0.0947, train acc: 0.3583, val loss: 0.1001, val acc : 0.3121
Epoch 11/20, train loss: 0.0986, train acc: 0.3668, val loss: 0.0846, val acc : 0.3774
Epoch 12/20, train loss: 0.0927, train acc: 0.3628, val loss: 0.0929, val acc : 0.3458
Epoch 13/20, train loss: 0.0931, train acc: 0.3605, val loss: 0.0930, val acc : 0.3446
Epoch 14/20, train loss: 0.0931, train acc: 0.3635, val loss: 0.0958, val acc : 0.3356
Epoch 15/20, train loss: 0.0937, train acc: 0.3659, val loss: 0.1190, val acc : 0.2893
Epoch 16/20, train loss: 0.0925, train acc: 0.3662, val loss: 0.0881, val acc : 0.3417
Epoch 17/20, train loss: 0.0927, train acc: 0.3684, val loss: 0.1048, val acc : 0.3247
Epoch 18/20, train loss: 0.0927, train acc: 0.3687, val loss: 0.0908, val acc : 0.3416
Epoch 19/20, train loss: 0.0929, train acc: 0.3686, val loss: 0.1023, val acc : 0.3099
Epoch 20/20, train loss: 0.0923, train acc: 0.3713, val loss: 0.0980, val acc : 0.3216

After 20 epochs of training, it achieves 32% of accuracy on test set.